

# AI Meets Query Optimization

Luming Sun @ DW&BI

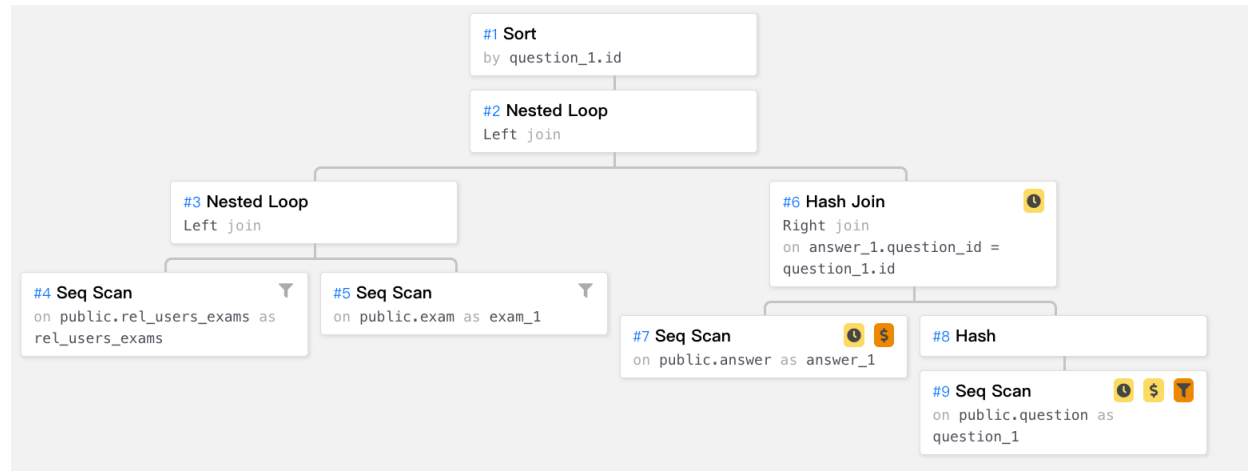
2022/9/19

# Query Optimization

Query optimizer is at the heart of the database systems

# Query Optimization

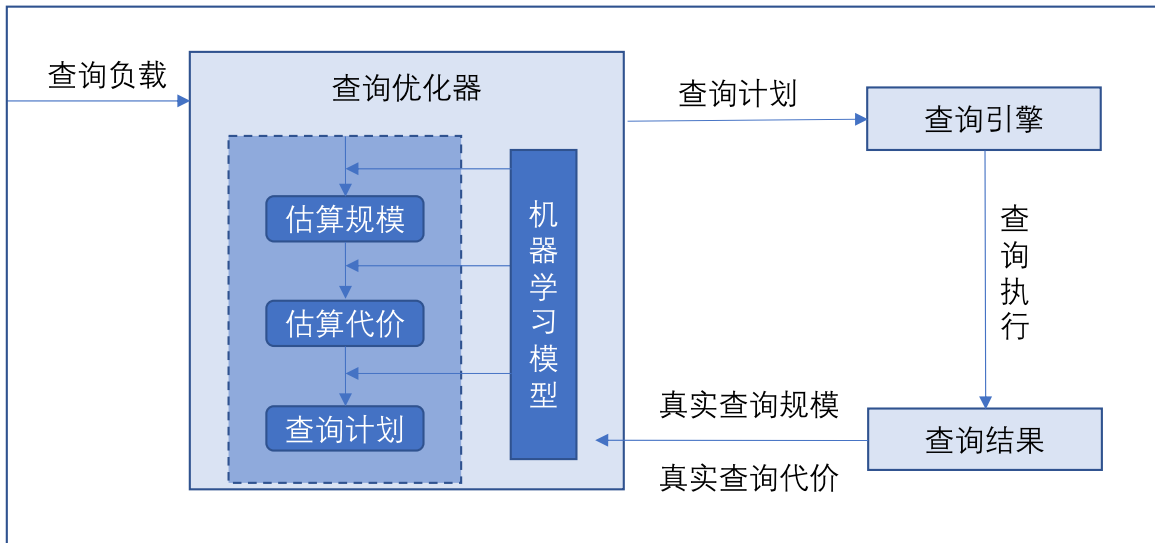
```
SELECT rel_users_exams.user_username AS rel_users_exams_user_username,  
       rel_users_exams.exam_id AS rel_users_exams_exam_id,  
       rel_users_exams.started_at AS rel_users_exams_started_at,  
       rel_users_exams.finished_at AS rel_users_exams_finished_at,  
       answer_1.id AS answer_1_id,  
       answer_1.text AS answer_1_text,  
       answer_1.correct AS answer_1_correct,  
       answer_1.fraction AS answer_1_fraction,  
       answer_1.question_id AS answer_1_question_id,  
       question_1.id AS question_1_id,  
       question_1.title AS question_1_title,  
       question_1.text AS question_1_text,  
       question_1.file AS question_1_file,  
       question_1.type AS question_1_type,  
       question_1.source AS question_1_source,  
       question_1.exam_id AS question_1_exam_id,  
       exam_1.id AS exam_1_id,  
       exam_1.title AS exam_1_title,  
       exam_1.date_from AS exam_1_date_from,  
       exam_1.date_to AS exam_1_date_to,  
       exam_1.created AS exam_1_created,  
       exam_1.created_by_ AS exam_1_created_by_,  
       exam_1.duration AS exam_1_duration,  
       exam_1.success_threshold AS exam_1_success_threshold,  
       exam_1.published AS exam_1_published  
FROM rel_users_exams LEFT OUTER  
JOIN exam AS exam_1  
  ON exam_1.id = rel_users_exams.exam_id LEFT OUTER  
JOIN question AS question_1  
  ON exam_1.id = question_1.exam_id LEFT OUTER  
JOIN answer AS answer_1  
  ON question_1.id = answer_1.question_id  
WHERE rel_users_exams.user_username = %(param_1)s  
      AND rel_users_exams.exam_id = %(param_2)s  
ORDER BY question_1.id;
```



# Cost-based optimizer

- A cost-based optimizer introduces a **plan enumeration algorithm** to find a (sub)plan, and then uses a cost model to obtain the cost of that plan, and selects the plan with the lowest cost.
- In the cost model, **cardinality**, the number of tuples through an operator, plays a crucial role.

# Why learning-based methods can help?



- Machine Learning can be used in:
  - Modeling data distribution & data correlation
  - Optimize function parameter
  - Markov decision process

# Cardinality/Selectivity Estimation

# What is Cardinality/Selectivity

```
: SELECT *  
FROM      WHERE      > 15  
AND      = 'Male';
```

$\text{Card}(Q) = 4$

$\text{Sel}(Q) = \text{Card}(Q) / \#row$   
 $= 4/9 = 0.444$

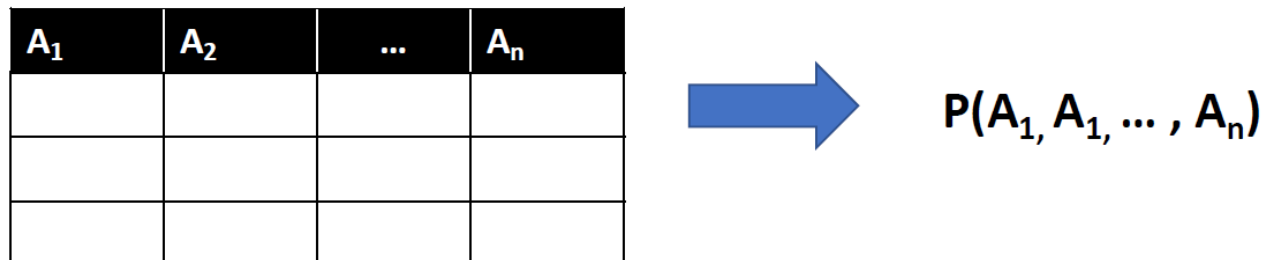
age	gender	GPA
21	Female	3.42
20	Male	2.58
18	Female	2.79
20	Female	3.98
24	Female	3.71
20	Male	3.50
21	Male	4.0
23	Female	3.66
22	Male	3.12

# How Learned Selectivity Estimators Work

- Methodology 1: **Query-driven**
  - Key Idea: Model as a Regression problem



- Methodology 2: **Data-driven**
  - Key Idea: Model as a Joint Distribution Estimation problem





# Methodology 1: Query-Driven

## Training

Query Pool

Q1: `SELECT * FROM Student WHERE age > 20;`  
Q2: `SELECT * FROM Student WHERE GPA < 3.5 AND GPA > 3.0;`  
Q3: `SELECT * FROM Student WHERE gender = 'Female';`  
...

Labels

4

2

5

...

Featurize

Q1: `<0.8, 1.0, 0.0, 0.0, 0.0, 1.0>` 4  
Q2: `<0.0, 1.0, 0.0, 1.0, 0.3, 0.6>` 2  
Q3: `<0.0, 1.0, 1.0, 1.0, 0.0, 1.0>` 5  
...

Train



Regression Model

## Inference

Q: `SELECT * FROM Student WHERE age > 15 AND gender = "Male"`

Featurize

Q: `<0.0, 0.9, 0.0, 1.0, 0.8, 1.0>`

Inference

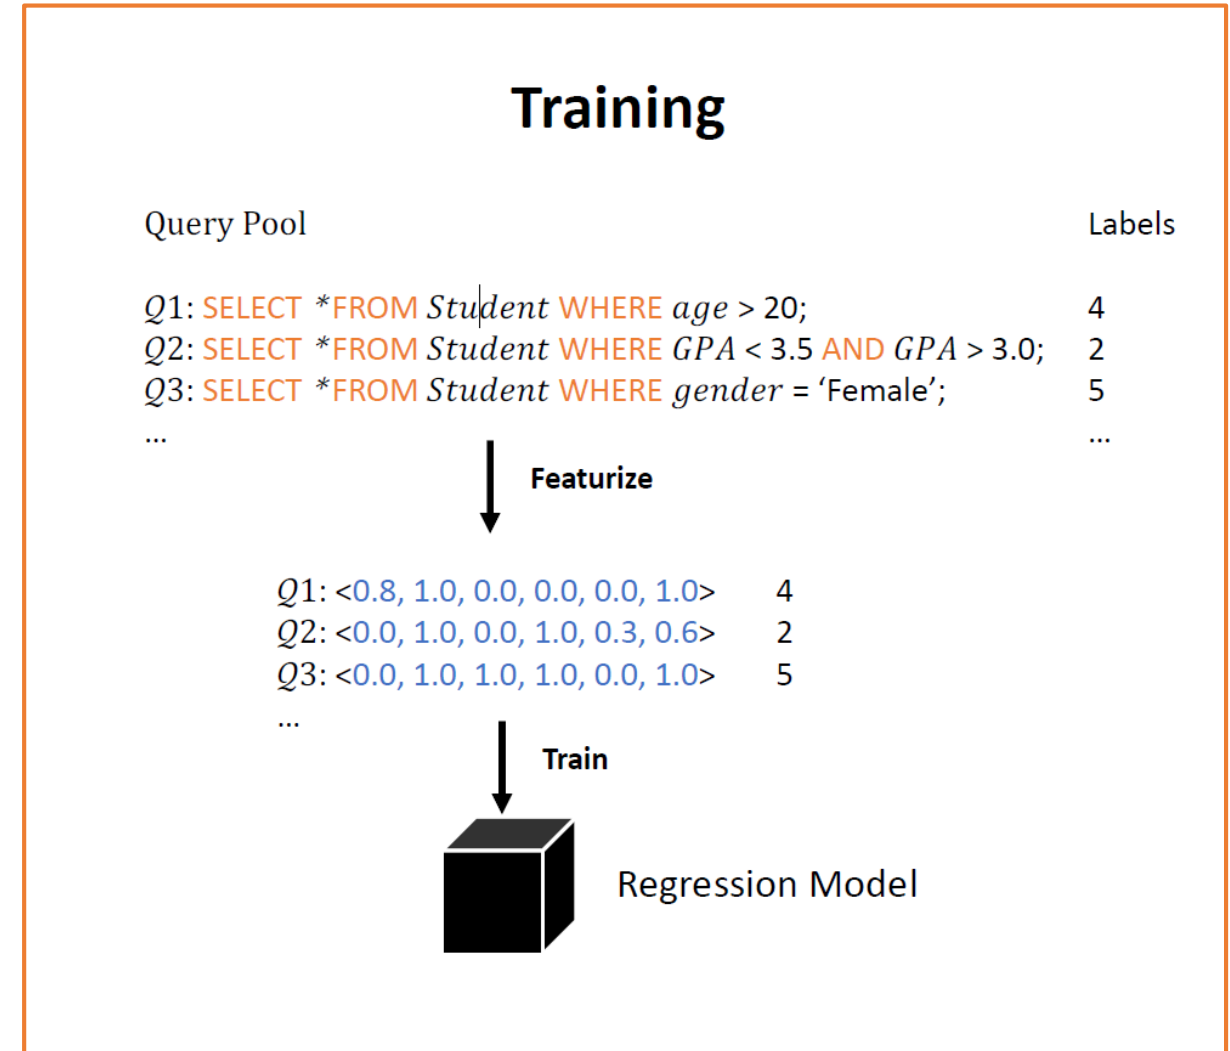


Estimation: 4!

# Methodology 1: Query-Driven

Single table:

- **LW-XGB** [Dutt, A et al. VLDB 19]  
XGBoost+ Heuristic statistics
- **LW-NN** [Dutt, A et al. VLDB 19]  
NN+ Heuristic statistics
- **QuickSel** [Yongjoo, P et al. SIGMOD 20]  
Mixture Model



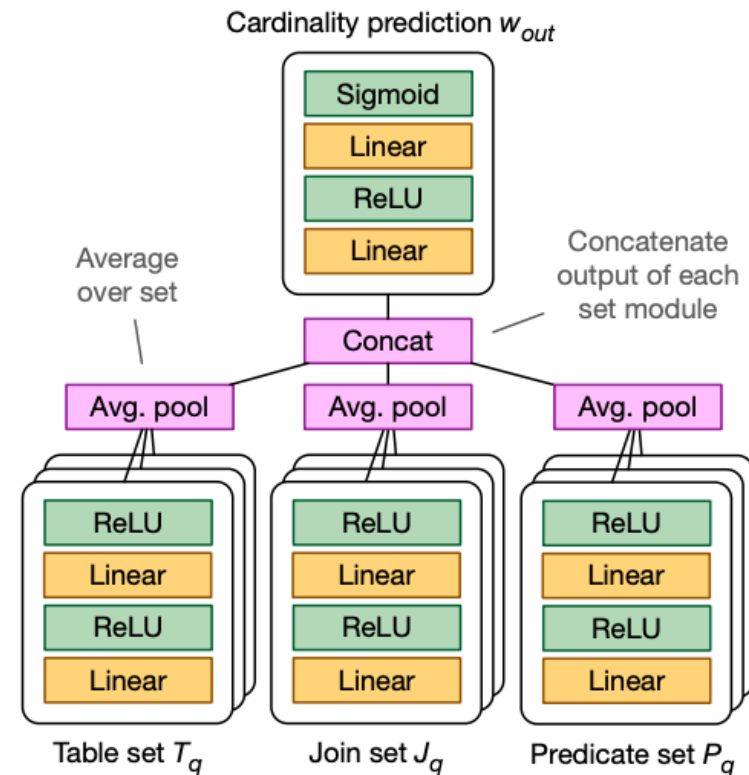
# Methodology 1: Query-Driven

Multiple tables:

- **LW-XGB** [Dutt, A et al. VLDB 20]  
XGBoost + incremental data gen
- **MSCN** [Kipf, A et al. CIDR 19]  
Neural Network + Sampling

```
SELECT COUNT(*) FROM title t, movie_companies mc WHERE t.id = mc.movie_id AND t.production_year > 2010 AND mc.company_id = 5
```

Table set  $\{ \underbrace{[0\ 1\ 0\ 1\ \dots\ 0]}_{\text{table id}}, \underbrace{[0\ 0\ 1\ 0\ \dots\ 1]}_{\text{samples}} \}$  Join set  $\{ \underbrace{[0\ 0\ 1\ 0]}_{\text{join id}} \}$  Predicate set  $\{ \underbrace{[1\ 0\ 0\ 0\ 0\ 1\ 0\ 0\ 0.72]}_{\text{column id value}}, \underbrace{[0\ 0\ 0\ 1\ 0\ 0\ 1\ 0\ 0.14]}_{\text{operator id}} \}$



# Query-Driven Shortcomings:

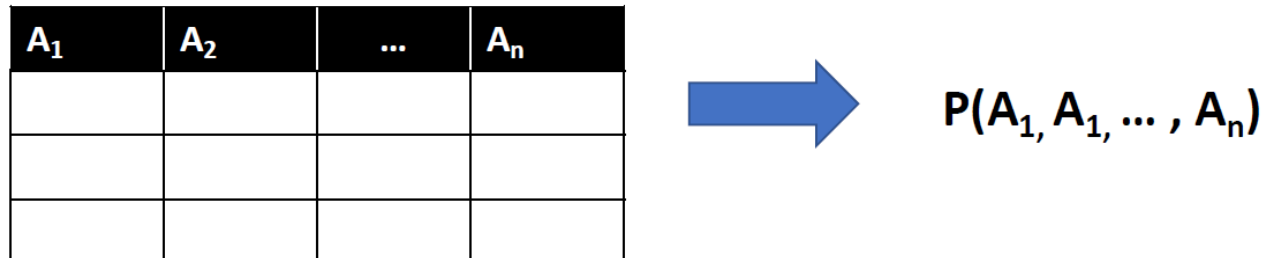
- Require large amount of training data
- Violate basic rule of selectivity estimation
  - - Monotonicity:  $sel(1 < x < 2) \leq sel(1 < x < 3)$
  - - Validity:  $sel(1 < x < 0) = 0$
  - - Consistency:  $sel(1 < x \leq 2) + sel(2 < x < 3)$ 
    - $= sel(1 < x < 3)$

# How Learned Selectivity Estimators Work

- Methodology 1: **Query-driven**
  - Key Idea: Model as a Regression problem



- Methodology 2: **Data-driven**
  - Key Idea: Model as a Joint Distribution Estimation problem



# Methodology 2: Data-Driven

## Training

age	gender	GPA
21	Female	3.42
20	Male	2.58
18	Female	2.79
20	Female	3.98
24	Female	3.71
20	Male	3.50
21	Male	4.0
23	Female	3.66
22	Male	3.12

Train



$P(\text{age}, \text{gender}, \text{GPA})$

Joint Distribution Estimation Model

## Inference

Q: `SELECT * FROM Student`  
`WHERE age > 15 AND gender = "Male"`



$P(\text{age} > 20, \text{gender} = \text{"Male"})$



Inference



Estimation: 4!

# Methodology 2: Data-Driven

- **Naru** [Yang, Z et al. VLDB 20]  
Auto-regressive Model
- **DeepDB** [Hilprecht, B et al. VLDB 20]  
Sum Product Network  
(Graphical Model)
- **FACE** [Wang, J et al. VLDB 2022]  
Normalizing Flow

## Training

age	gender	GPA
21	Female	3.42
20	Male	2.58
18	Female	2.79
20	Female	3.98
24	Female	3.71
20	Male	3.50
21	Male	4.0
23	Female	3.66
22	Male	3.12

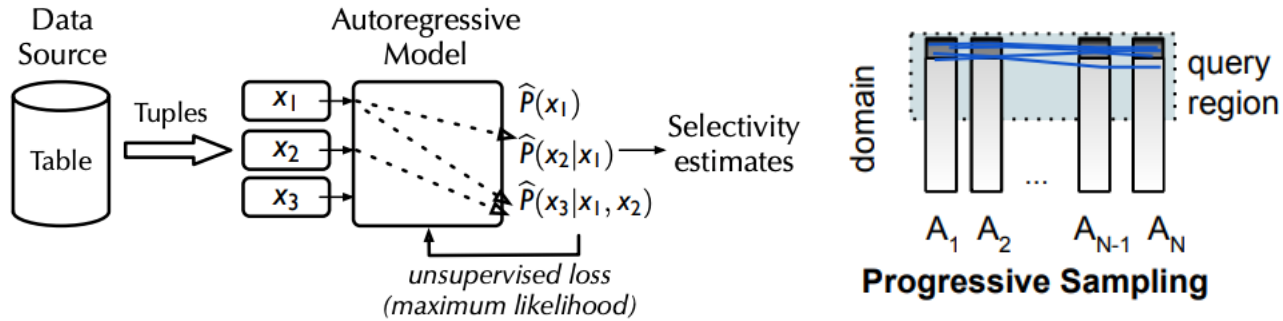
Train



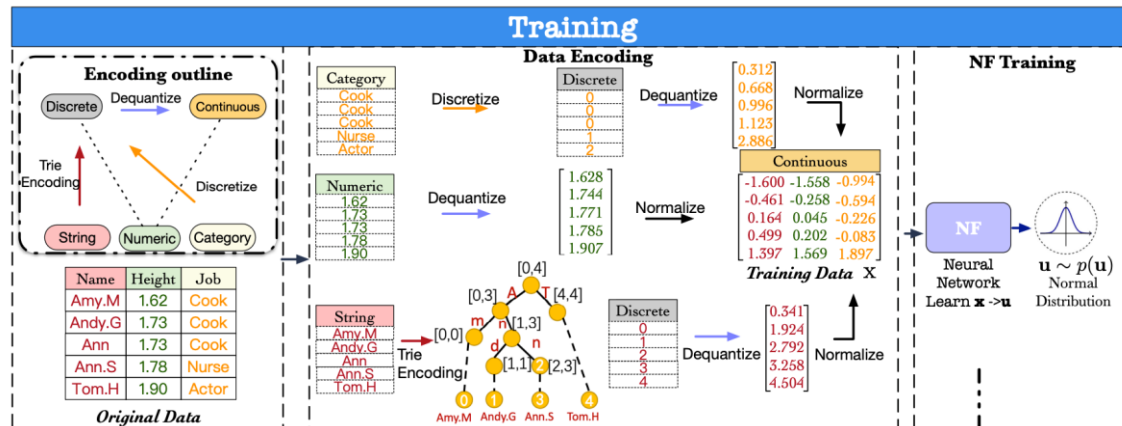
$P(\text{age, gender, GPA})$

Joint Distribution Estimation Model

# Data-Driven Methods



Naru: Auto-regressive model



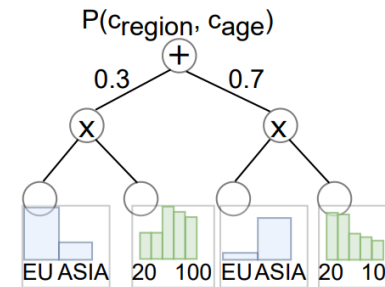
FACE: Normalizing flow

c_id	c_age	c_region
1	80	EUROPE
2	70	EUROPE
3	60	ASIA
4	20	EUROPE
...	...	...
998	20	ASIA
998	25	EUROPE
999	30	ASIA
1000	70	ASIA

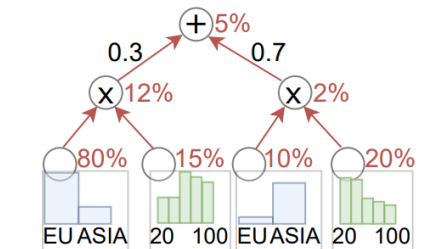
(a) Example Table

c_age	c_region
80	EUROPE
70	EUROPE
60	ASIA
20	EUROPE
...	...
...	...
20	ASIA
25	EUROPE
30	ASIA
70	ASIA

(b) Learning with Row/Column Clustering



(c) Resulting SPN



(d) Probability of European Customers younger than 30

DeepDB: Sum-product network



# Shortcomings

- Heavy costs on model training and inference
- Violate basic rule of selectivity estimation
  - - Monotonicity:  $sel(1 < x < 2) \leq sel(1 < x < 3)$
  - - Validity:  $sel(1 < x < 0) = 0$
  - - Consistency:  $sel(1 < x \leq 2) + sel(2 < x < 3)$ 
    - $= sel(1 < x < 3)$
  - - **Stability**:  $sel(1 < x < 2) = sel(1 < x < 2)$

# MOSE: A Monotonic Selectivity Estimator Using Learned CDF

- **Aim**

Reliable, accurate and efficient learned selectivity estimator

- **Problem Settings**

Multi-dimensional predicates on single table

- **Methodology**

Query-based

- **Key observation**

The joint cumulative distribution function (CDF) of the data in a table can be used to compute the selectivity for query range predicates

# CDF to Selectivity

- Multi-dimensional CDF:

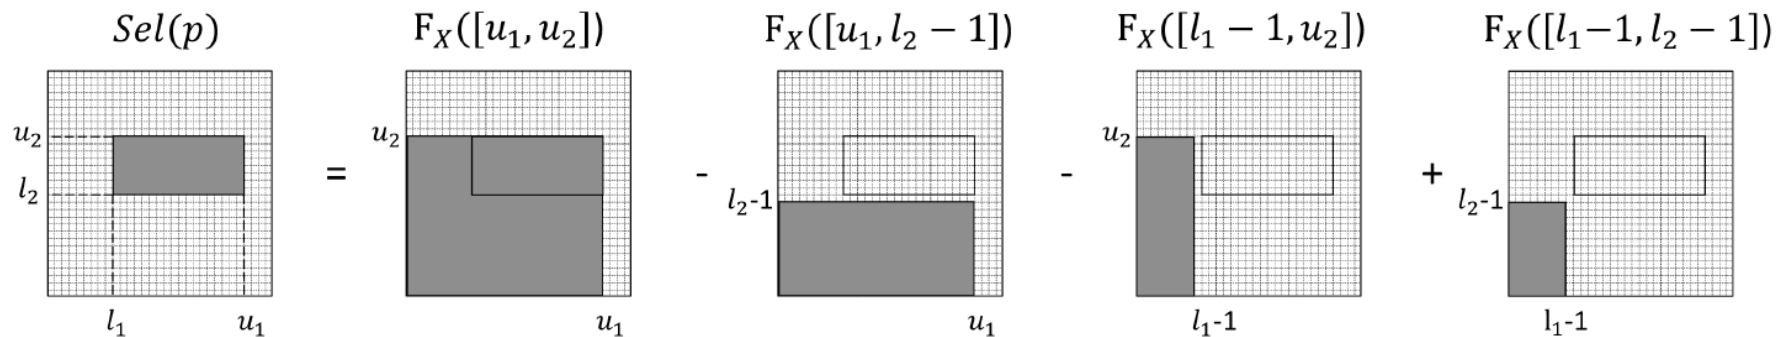
Random variable  $X = (X_1, X_2, \dots, X_d)$ ,

CDF:  $F_X(x) = \Pr(X_1 \leq x_1, X_2 \leq x_2, \dots, X_d \leq x_d)$

- CDF to selectivity

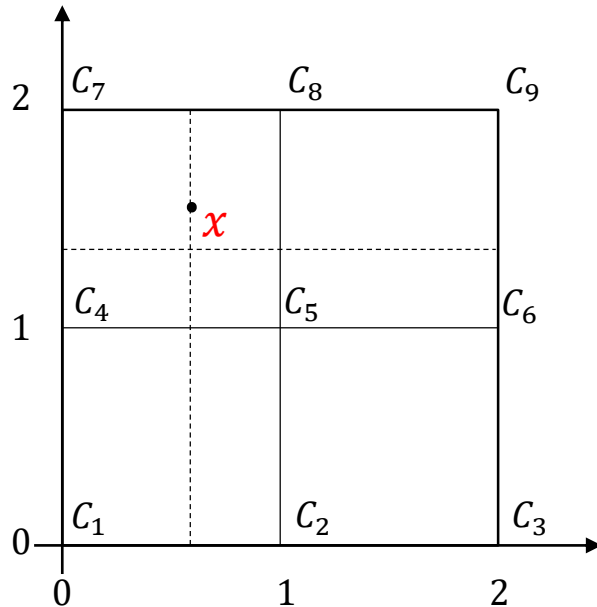
$$sel(p) = \sum_{\forall x_i \in \{l_i-1, u_i\}} \left\{ \left( \prod_{i=1}^d s(i) \right) F_X(x) \right\}$$

$$sel(p) = F_X([u_1, u_2]) - F_X([u_1, l_2 - 1]) - F_X([l_1 - 1, u_2]) + F_X([l_1 - 1, l_2 - 1])$$



# CDF Learner

- Monotonic Lattice Regression Model



C: Lattice Node

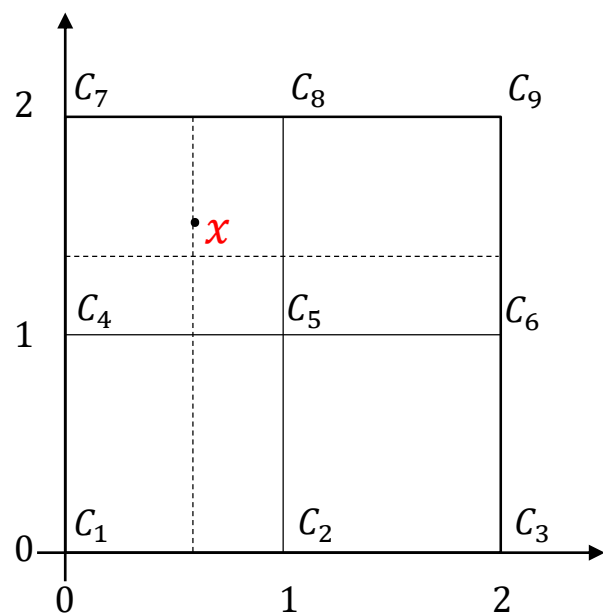
Theta: Lattice Parameter

$$\hat{y} = F_X(x) = \sum_{j=1}^m \phi(x)_j \theta_j, \quad \sum_{j=1}^m \phi(x)_j C_j = x, \quad \sum_{j=1}^m \phi(x)_j = 1.$$

$$\begin{aligned} \theta &= \arg \min_{\theta \in \mathbb{R}^m} \sum_{i=1}^n (\hat{y}_i - y_i)^2 \\ &= \arg \min_{\theta \in \mathbb{R}^m} \sum_{i=1}^n \left( \left( \sum_{j=1}^m \phi(x)_{ij} \theta_j \right) - y_i \right)^2. \end{aligned}$$

$$\theta = \arg \min_{\theta \in \mathbb{R}^m} \sum_{i=1}^n (\hat{y}_i - y_i)^2, \text{ s.t. } A\theta^T \leq 0.$$

# Monotonic Constrain



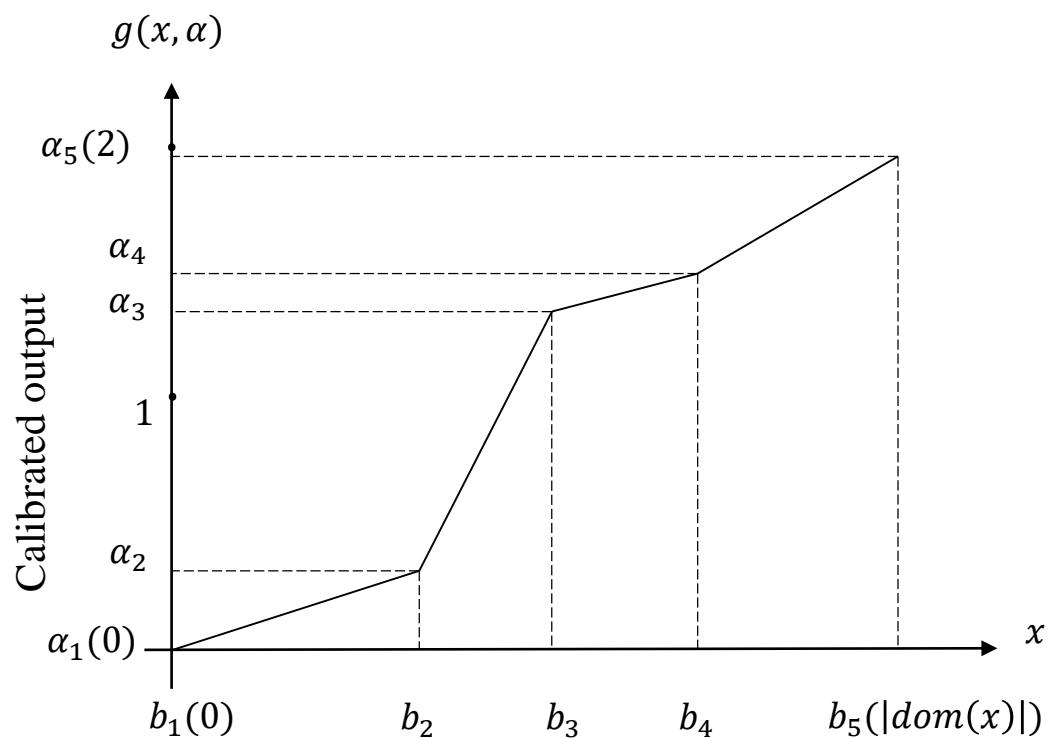
C: Lattice Node

Theta: Lattice Parameter

$$\theta = \arg \min_{\theta \in \mathbb{R}^m} \sum_{i=1}^n (\hat{y}_i - y_i)^2, s.t. A\theta^T \leq 0.$$

$$\begin{bmatrix} 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & -1 \\ 1 & 0 & 0 & -1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & -1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & -1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & -1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & -1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & -1 \end{bmatrix} \begin{bmatrix} \theta_1 \\ \theta_2 \\ \theta_3 \\ \theta_4 \\ \theta_5 \\ \theta_6 \\ \theta_7 \\ \theta_8 \\ \theta_9 \end{bmatrix} \leq 0 \Leftrightarrow \begin{array}{l} \theta_1 \leq \theta_2 \\ \theta_2 \leq \theta_3 \\ \theta_4 \leq \theta_5 \\ \theta_5 \leq \theta_6 \\ \theta_7 \leq \theta_8 \\ \theta_8 \leq \theta_9 \\ \theta_1 \leq \theta_4 \\ \theta_2 \leq \theta_5 \\ \theta_3 \leq \theta_6 \\ \theta_4 \leq \theta_7 \\ \theta_5 \leq \theta_8 \\ \theta_6 \leq \theta_9 \end{array}$$

# Attribute-Aware Calibration



$$\theta, \alpha = \arg \min_{\theta, \alpha} \sum_{i=1}^n \left( \left( \sum_{j=1}^m \phi(g(x, \alpha))_{ij} \theta_j \right) - y_i \right)^2 + \lambda R(\theta)$$

$$s.t. \quad A\theta^T \leq 0 \text{ and } B\alpha^T \leq 0,$$

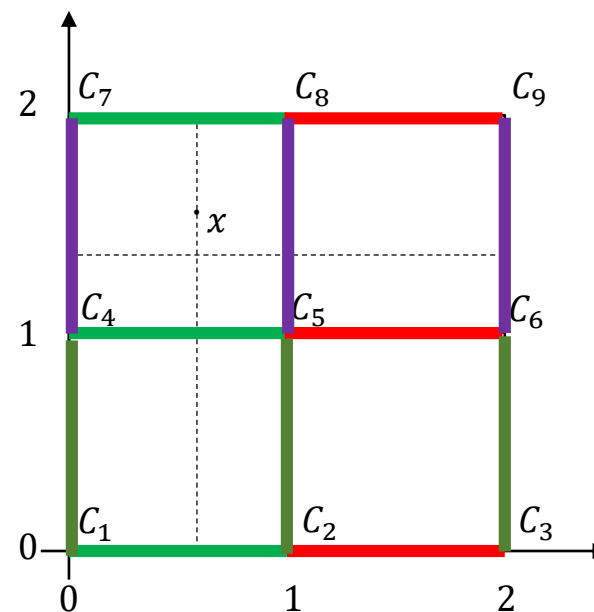
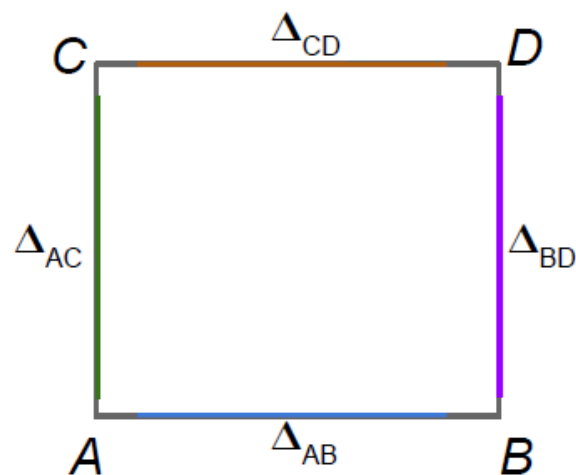
# Cell-Wise Regularizer

**Graph Laplacian:**  
flatter function

**Penalizes:**

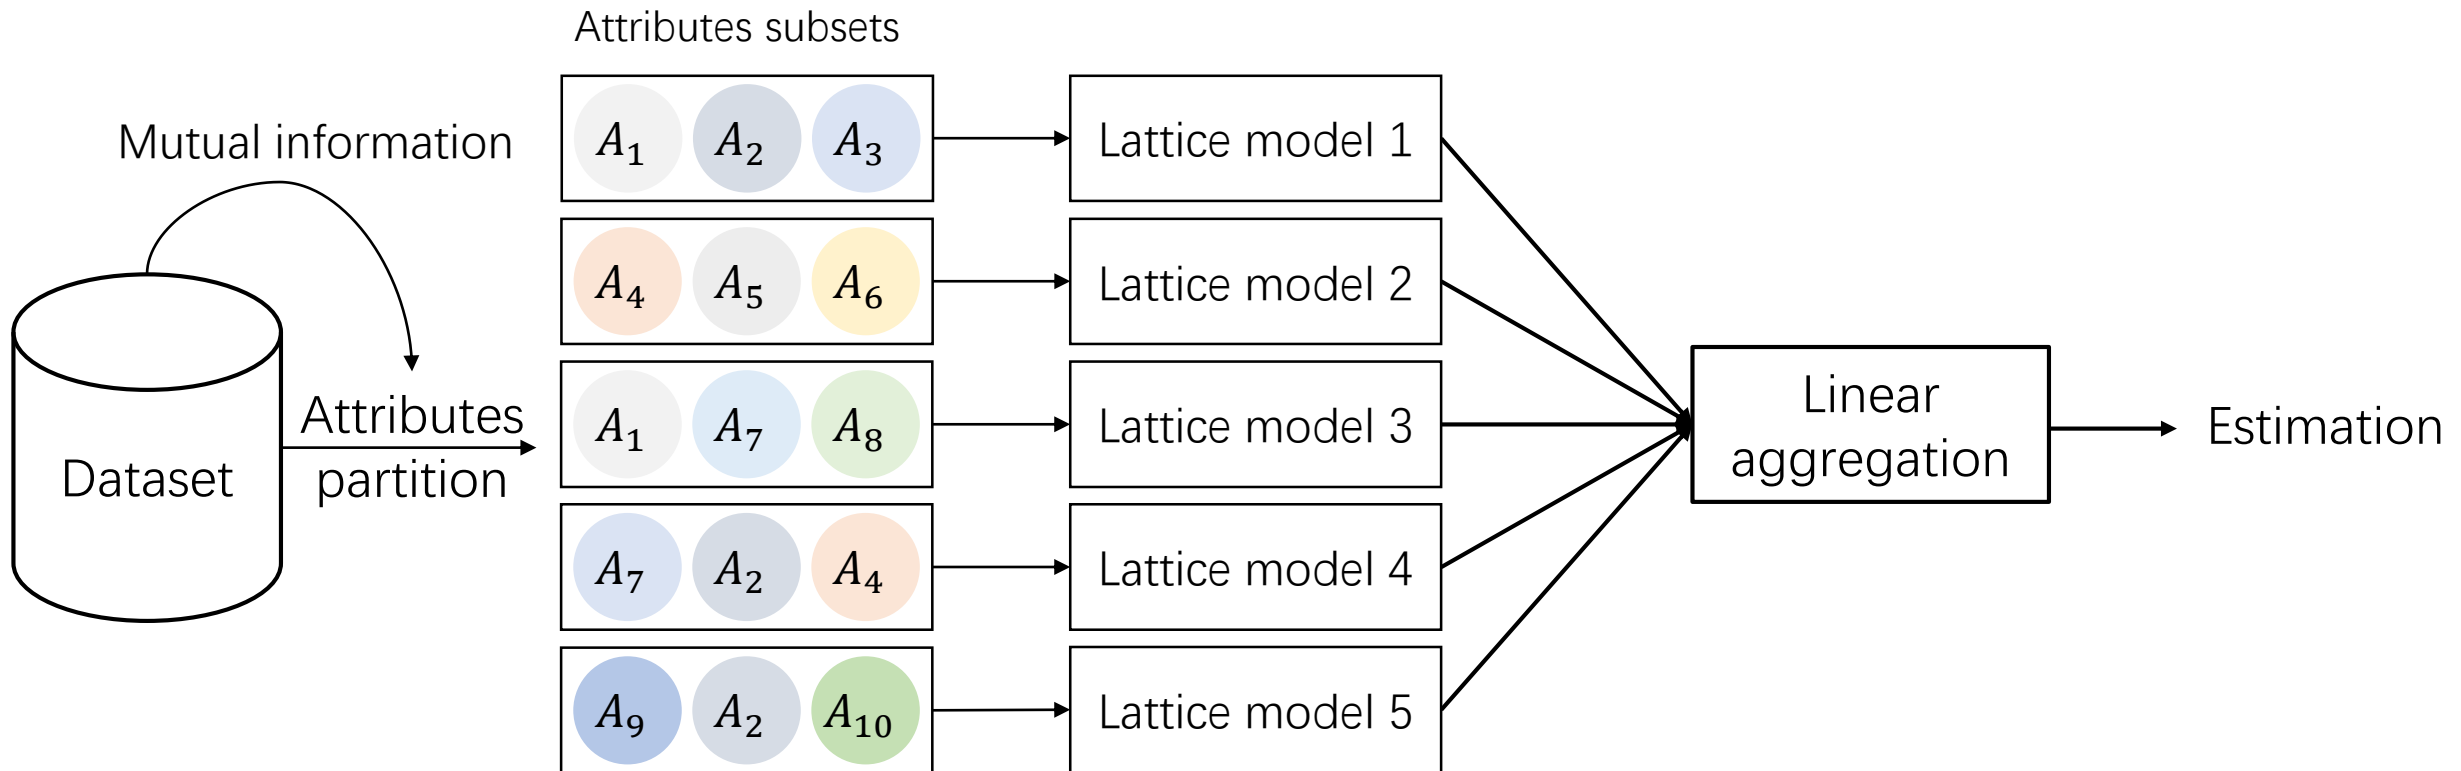
$$(A-C)^2 + (A-B)^2 + (C-D)^2 + (B-D)^2$$

$$= \Delta_{AC}^2 + \Delta_{AB}^2 + \Delta_{CD}^2 + \Delta_{BD}^2$$



$$R(\theta) = \sum_{i=1}^d \sum_{\substack{C_r, C_s \text{ such that} \\ C_r \text{ and } C_s \text{ adjacent on dimension } i}} H_i(C_r, C_s) (\theta_r - \theta_s)^2$$

# Lattice Ensemble





# ACTIVE DATA GENERATOR

- Challenges:
  - Infinite query space (NOT pool based active learning)
  - Regression problem (NO model uncertainty)
- Solution
  - Picking the lattice cells that are most valuable or necessary to optimize
  - Two factors: (1) cell accuracy; (2) cell density
  - Weighted lattice sampling

# Weighted Lattice Sampling

---

**Algorithm 2:** Active Data Generator

---

**input :**  $X_L$  is the initial labeled data,  
 $\theta$  is model weight of CDF learner,  
 $\mathcal{T}$  is the cost threshold of data collection,  
 $\epsilon$  is the cost function to label a data instance,  
 $B$  is the number of data selected in one batch,  
 $\mathcal{P}$  is a function to calculate cell sampling weight

**output:**  $X$  is the labeled training data

```
1  $X \leftarrow X_L$  // total training set
2  $t \leftarrow 0$  // initialize total cost
3 while  $t < \mathcal{T}$  do
4    $\theta \leftarrow \text{TrainModelWith}(X)$ 
5    $\text{Error} \leftarrow \text{Evaluate}(\theta, X)$ 
6    $P_{\text{TopError}} \leftarrow \text{Top}(X, \text{Error}, K)$ 
7    $\mathcal{P}_c \leftarrow X, P_{\text{TopError}}$ 
8    $\text{cells} \leftarrow \text{WeightedSampling}(\mathcal{P}_c, B)$ 
9    $X_A \leftarrow \text{RandomPointGenerate}(\text{cells})$ 
10   $X_{AL} \leftarrow \text{ExecuteQuery}(X_A)$ 
11   $t = t + \epsilon(X_A)$ 
12   $X = X \cup X_{AL}$ 
13 return  $X$ 
```

---

$$\mathcal{P}_c = \frac{1 + \omega k_c}{1 + M_c},$$

$M_c$ : points in cell C

$k_c$ : k points of C are in the  
TOP-K worst estimation

# Accuracy

TABLE 2: Selectivity estimation accuracy on DMV

Estimator	Training data size				
	200	400	600	800	1000
LWM	0.03474	0.02576	0.01817	0.01758	0.01607
NN	0.04787	0.03082	0.02153	0.01803	0.01577
QuickSel	0.02151	0.01421	0.01296	0.01125	0.01027
MOSE	0.00674	0.00543	0.00463	0.00429	0.00393

TABLE 3: Selectivity estimation accuracy on Forest

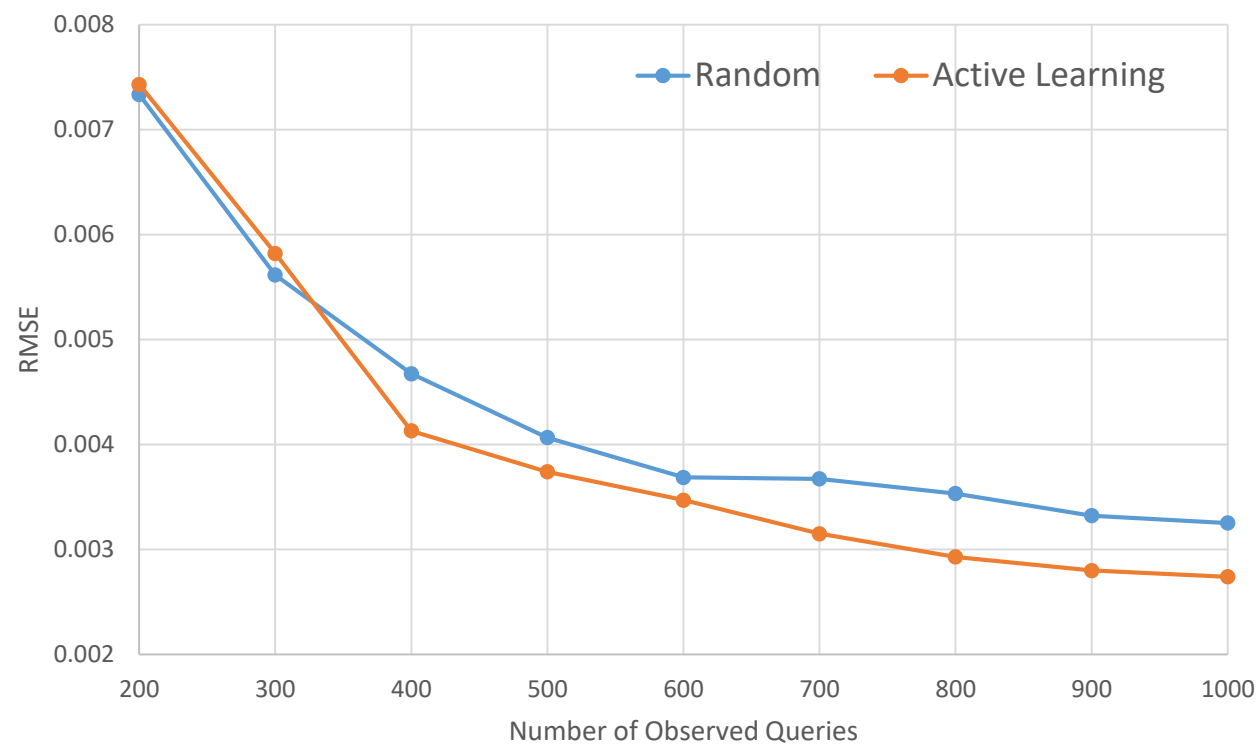
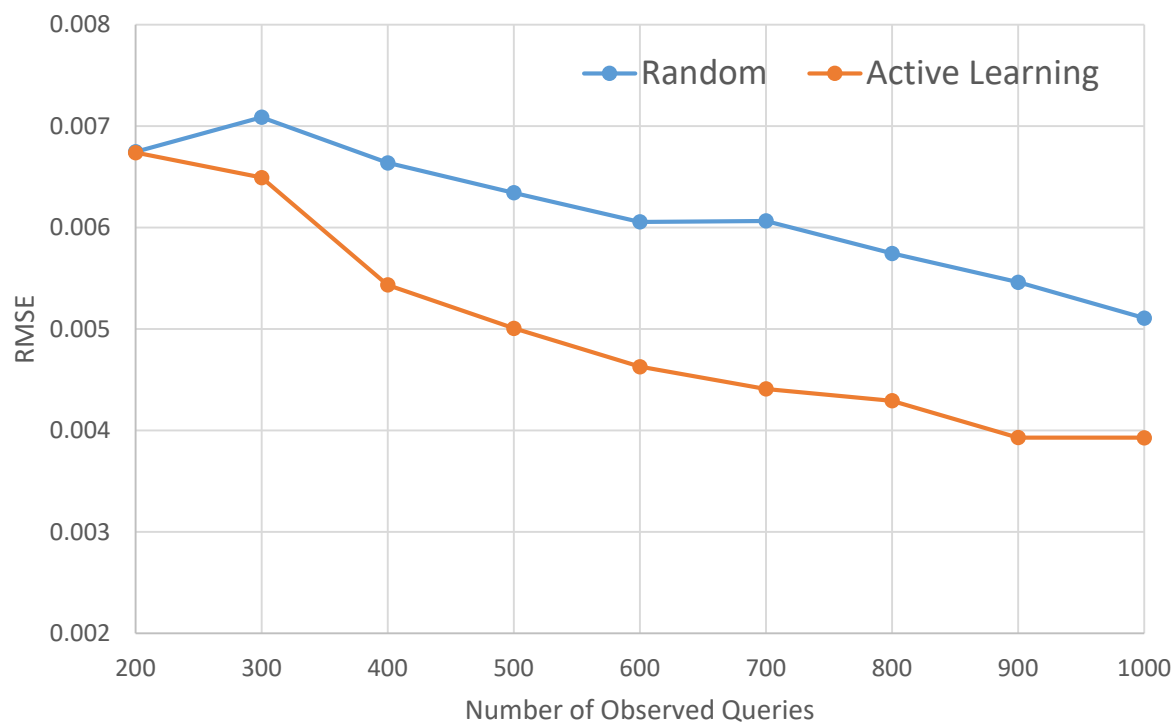
Estimator	Range predicates dimension				
	2D	4D	6D	8D	10D
AVI	0.23020	0.06069	0.01060	0.00240	0.000582
Sampling	0.00642	0.01164	0.00452	0.00946	0.000718
Naru	0.20113	0.59320	0.56103	0.10131	0.308497
LWM	0.03125	0.01573	0.00729	0.00229	0.000574
NN	0.00638	0.01226	0.00943	0.00240	0.000582
QuickSel	0.00470	0.00773	0.00382	0.83949	0.000590
MOSE	0.00419	0.00544	0.00274	0.00223	0.000555

# Ablation Experiments

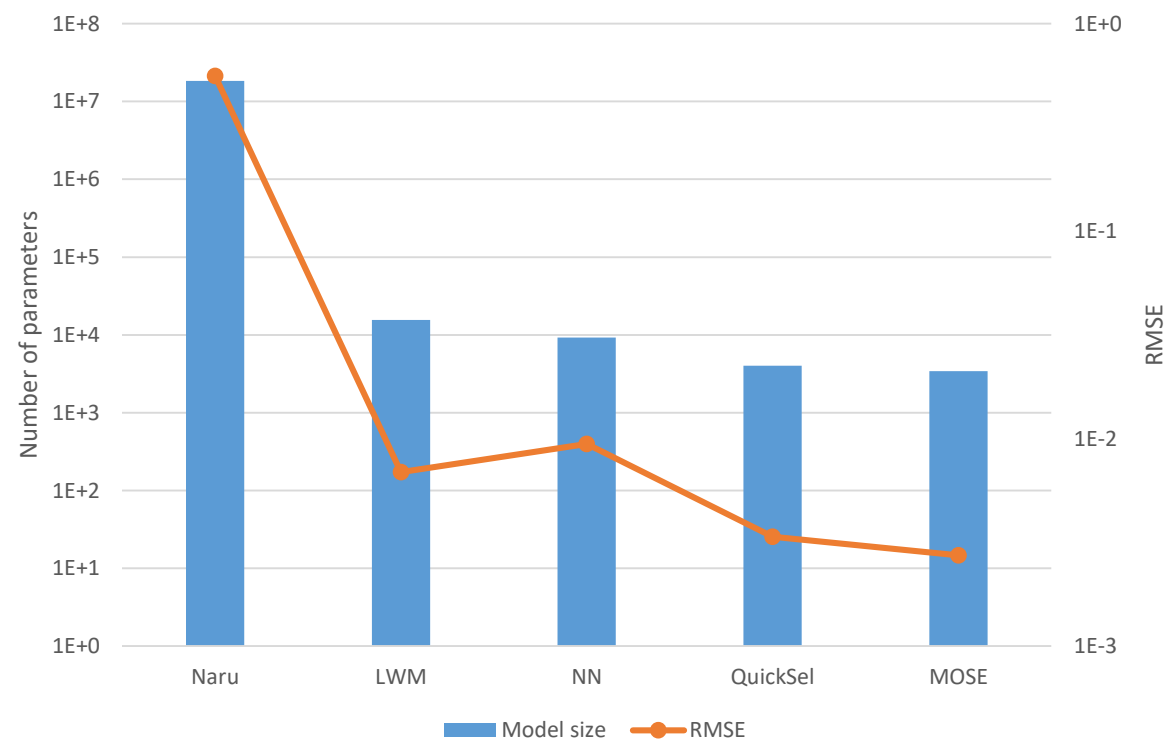
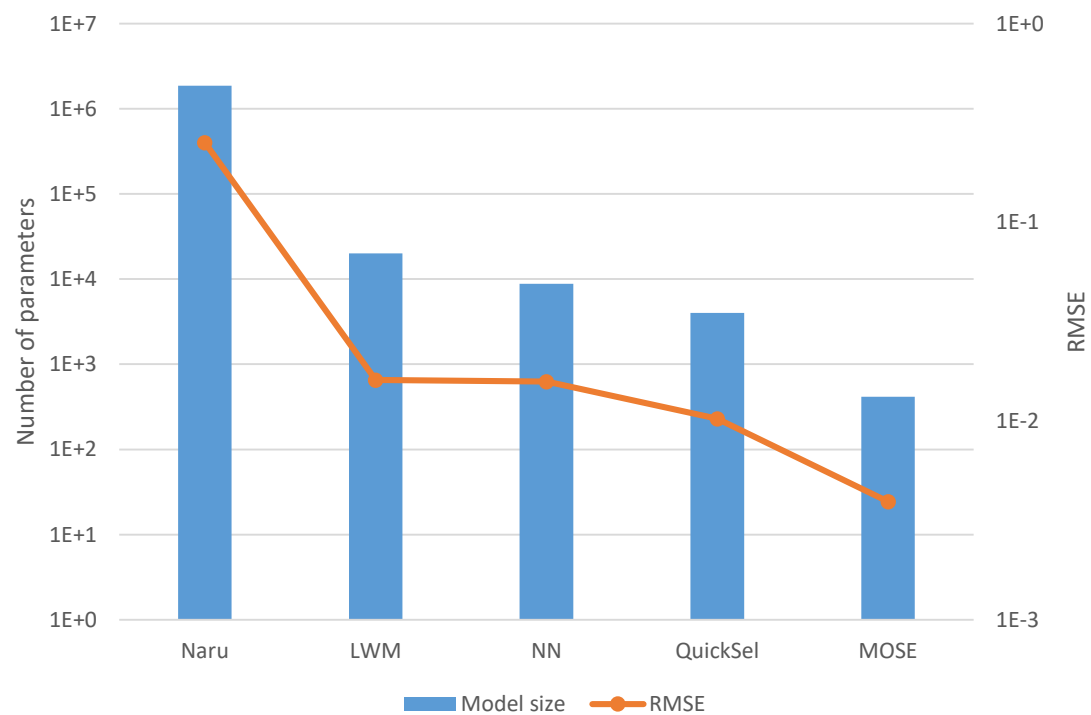
TABLE 4: Combination of calibration and regularizer

Combination method	RMSE
Laplacian regularizer + Uniform calibration	0.00713
Laplacian regularizer + A-A calibration	0.00540
C-W regularizer + Uniform calibration	0.00530
C-W regularizer + A-A calibration	0.00393

# Active Learning



# Model Size



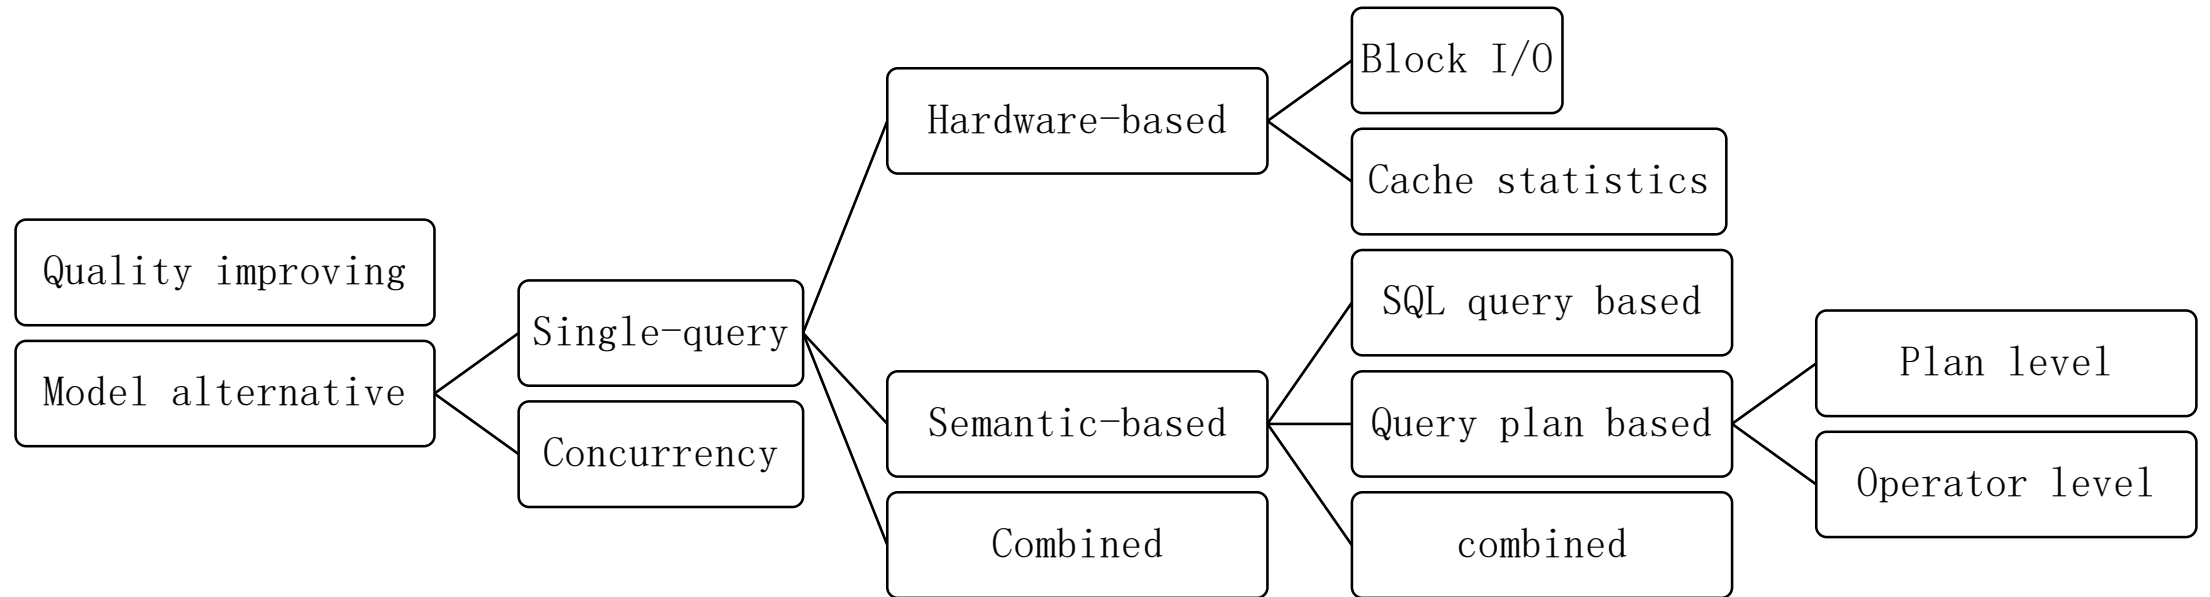
# Takeaways

- Learning-based methods can be accurate, efficient, lightweight
- But
  - Have longer-training time
  - May fail to catch up fast data change
  - Hard to trust
  - Are not practical in production

# Cost Estimation



# Learning-Based Cost Estimation



# Cost model in PostgreSQL

1. *seq\_page\_cost*( $c_s$ ): the cost of a sequential disk page fetch
2. *random\_page\_cost* ( $c_r$ ): the cost of fetching a disk page randomly
3. *cpu\_tuple\_cost* ( $c_t$ ): the cost of processing a tuple
4. *cpu\_index\_tuple\_cost*( $c_i$ ): the cost of processing an index entry during an index scan
5. *cpu\_operator\_cost*( $c_o$ ): the cost of performing an operation

The cost of an operator,  $C_o$  is calculated as [8]:

$$C_o = n_s \cdot c_s + n_r \cdot c_r + n_t \cdot c_t + n_i \cdot c_i + n_o \cdot c_o \quad (2.1)$$

where

$n_s$ : number of disk pages fetched sequentially

$n_r$ : number of disk pages fetched randomly

$n_t$ : number of tuples processed

$n_i$ : number of index entries processed during an index scan

$n_o$ : number of operations performed



# Plan-structured cost learner

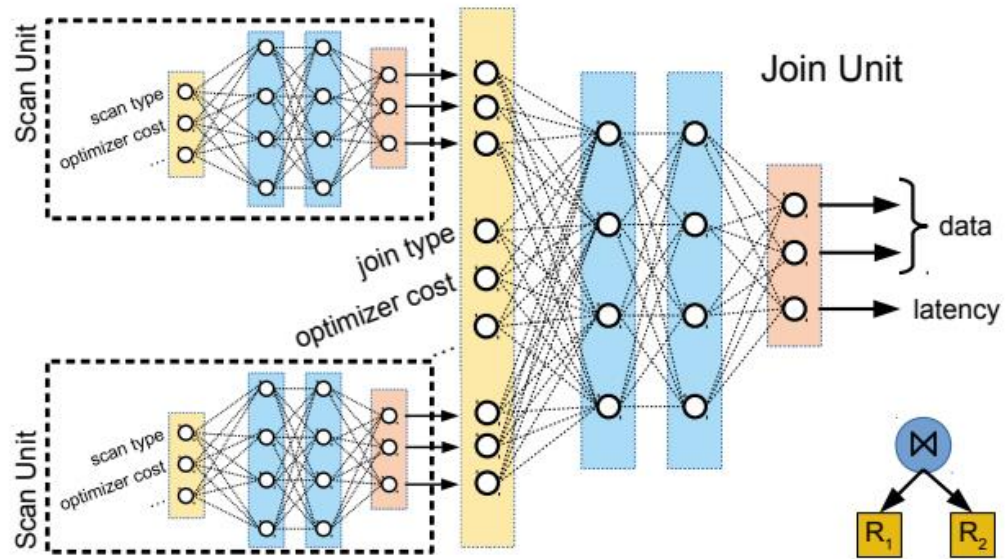
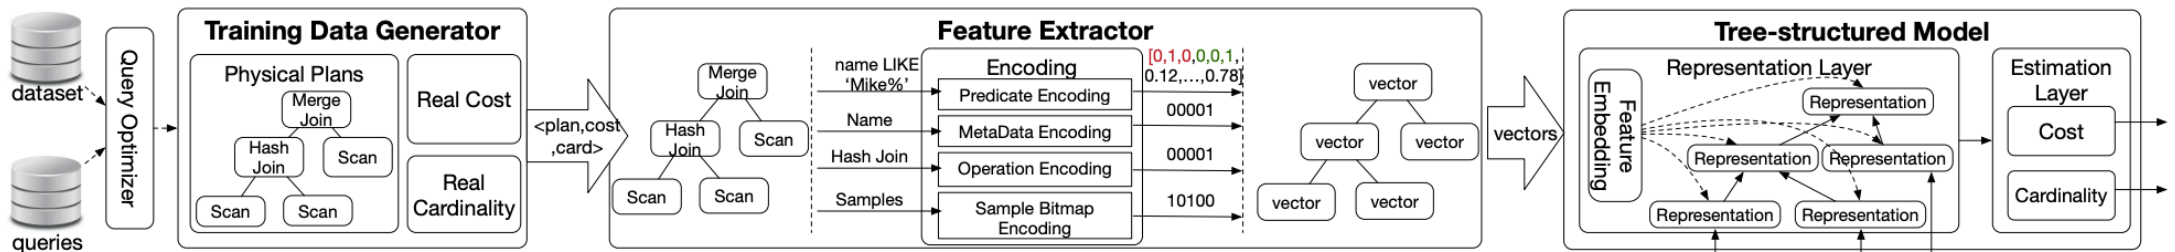


Figure 3: A neural network for a simple join query

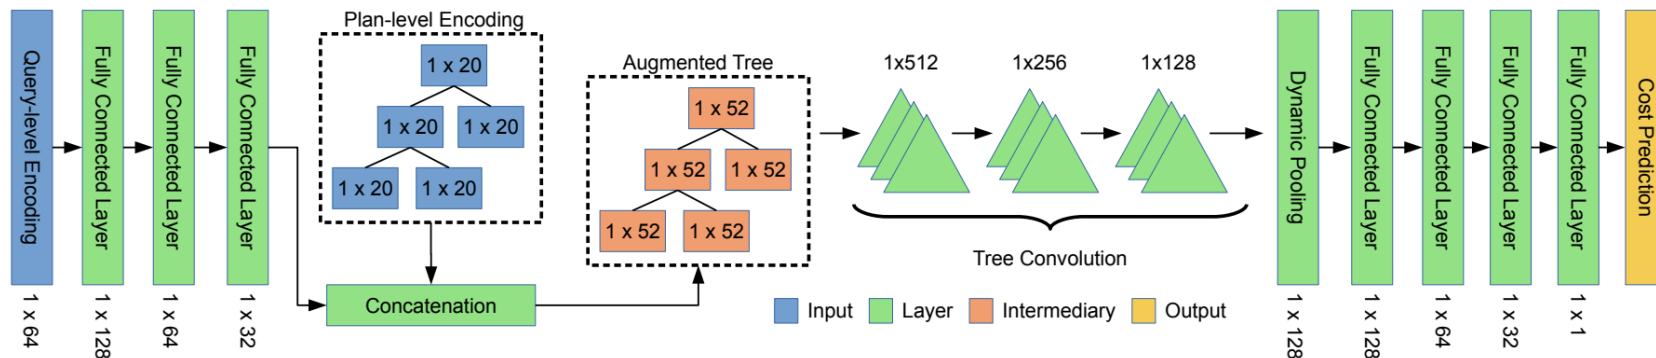
Plan-Structured Deep Neural Network Models for Query Performance Prediction (VLDB 2019)

# Plan-structured cost learner



**Figure 2:** Architecture of learning-based cost estimator

An End-to-End Learning-based Cost Estimator (VLDB 2019)



**Figure 5:** Value network architecture

Neo: A Learned Query Optimizer (VLDB 2019)

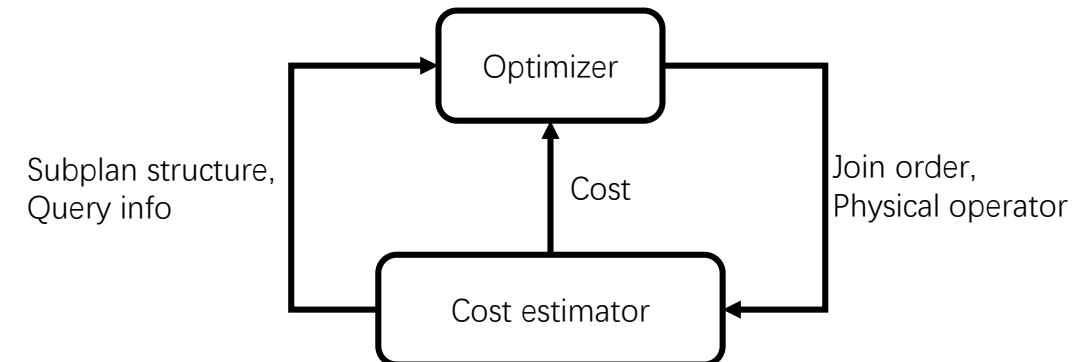
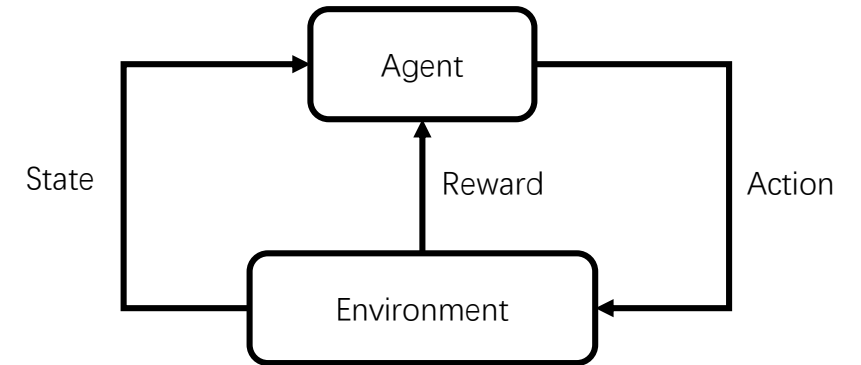
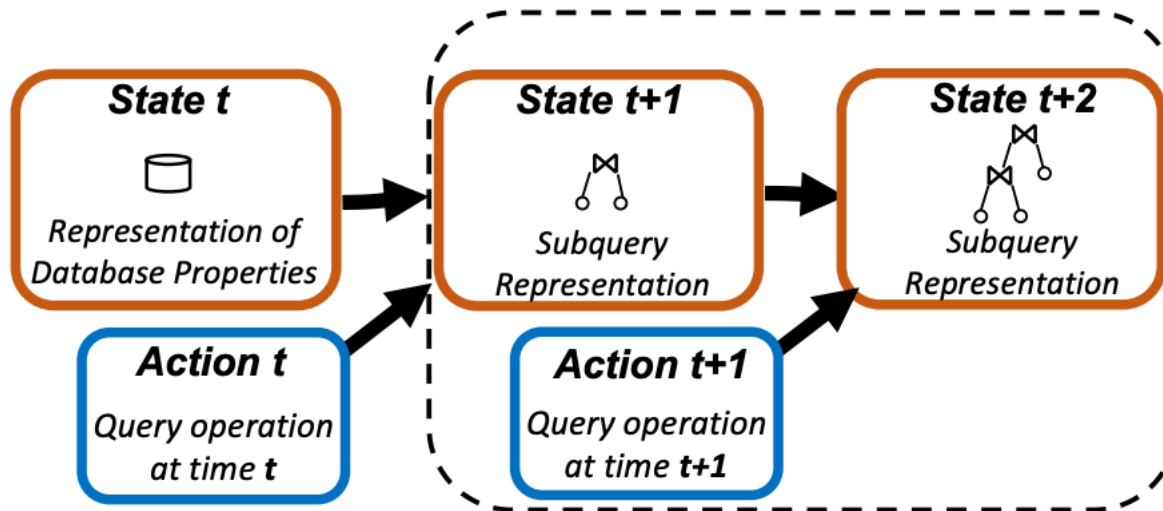
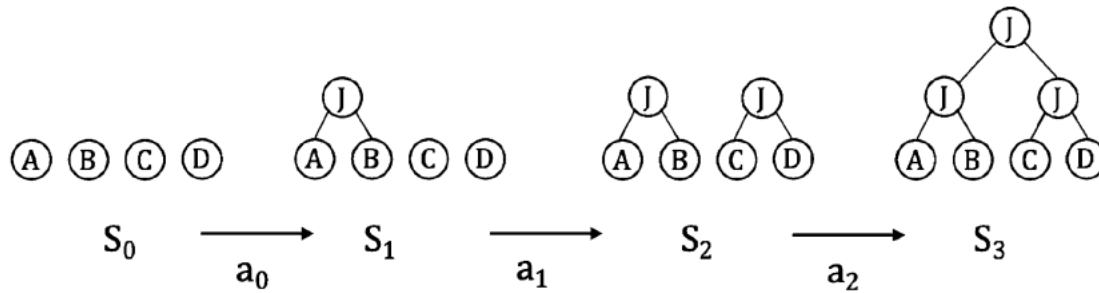
# Takeaways

- Learning-based cost estimation are accurate but not so efficient.
- Some learning-based cost estimation methods are used for scheduling not query optimization.
- Learning-based cost estimation (especially tree-structured) are not so practical in traditional cost-based query optimization.

Plan enumeration

# Query Plan Optimization

SELECT \* FROM A, B, C, D WHERE A.a1 = B.b1  
AND B.b1 = C.c1 AND C.c1 = D.d1;



# Query Plan Optimization

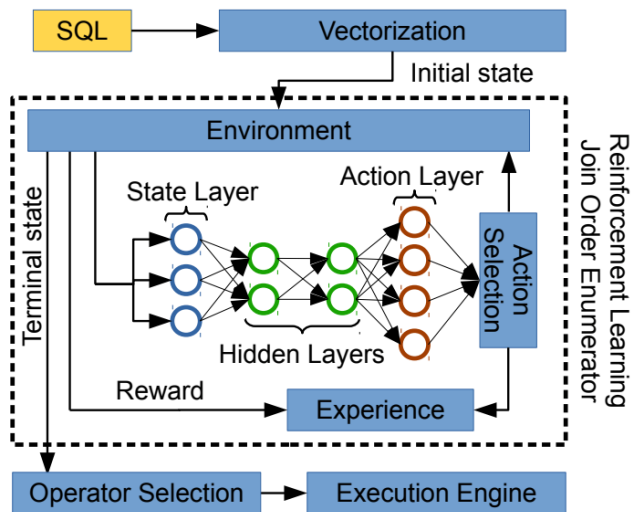


Figure 2: The ReJOIN Framework

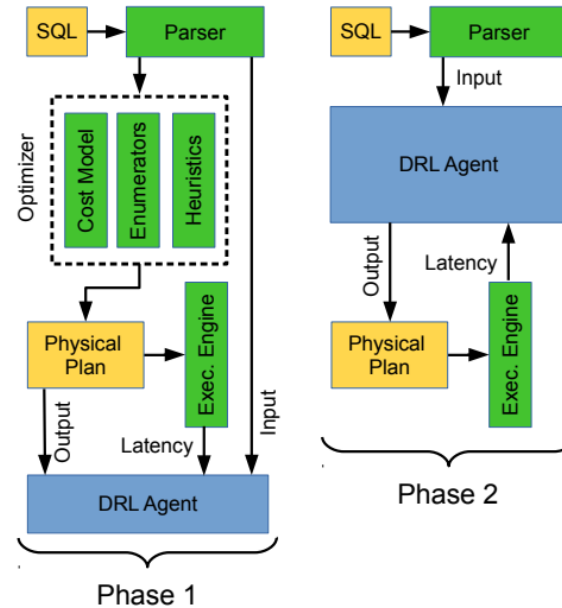


Figure 4: Learning from demonstration

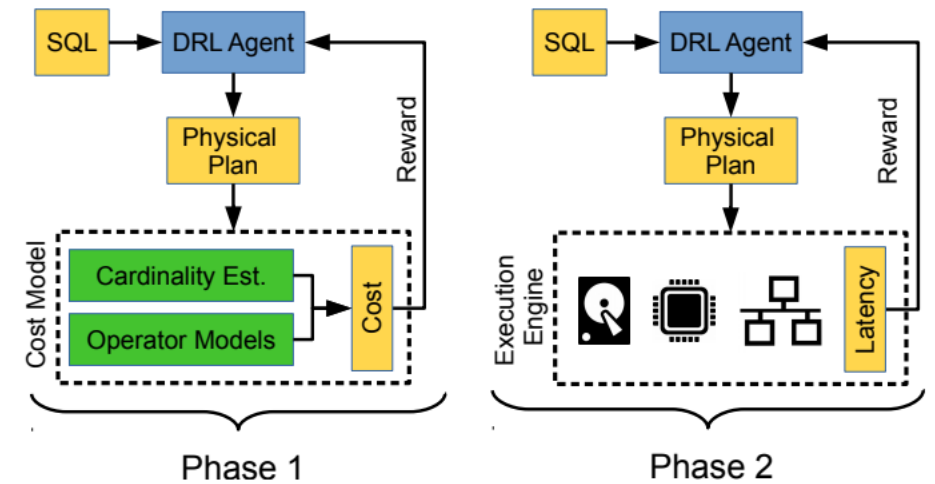


Figure 5: Cost Model Bootstrapping

1. Krishnan, Sanjay, et al. "Learning to optimize join queries with deep reinforcement learning." arXiv preprint arXiv:1808.03196 (2018).
2. Marcus, Ryan, and Olga Papaemmanouil. "Deep reinforcement learning for join order enumeration." Proceedings of the First International Workshop on Exploiting Artificial Intelligence Techniques for Data Management (2018): 3.
3. Ortiz, Jennifer, et al. "Learning state representations for query optimization with deep reinforcement learning." arXiv preprint arXiv:1803.08604 (2018).
4. Marcus, Ryan and Olga Papaemmanouil. "Towards a hands-free query optimizer through deep learning. " (CIDR 2019).
5. Marcus, Ryan, et al. "Neo: A learned query optimizer." (VLDB 2019).



# Shortcomings

- Long training time
- Inability to adjust to data and workload changes
- Tail catastrophe
- Black-box decisions
- Integration cost

# Bao: Making Learned query Optimization Practical (SIGMOD 2021 **Best Paper**)

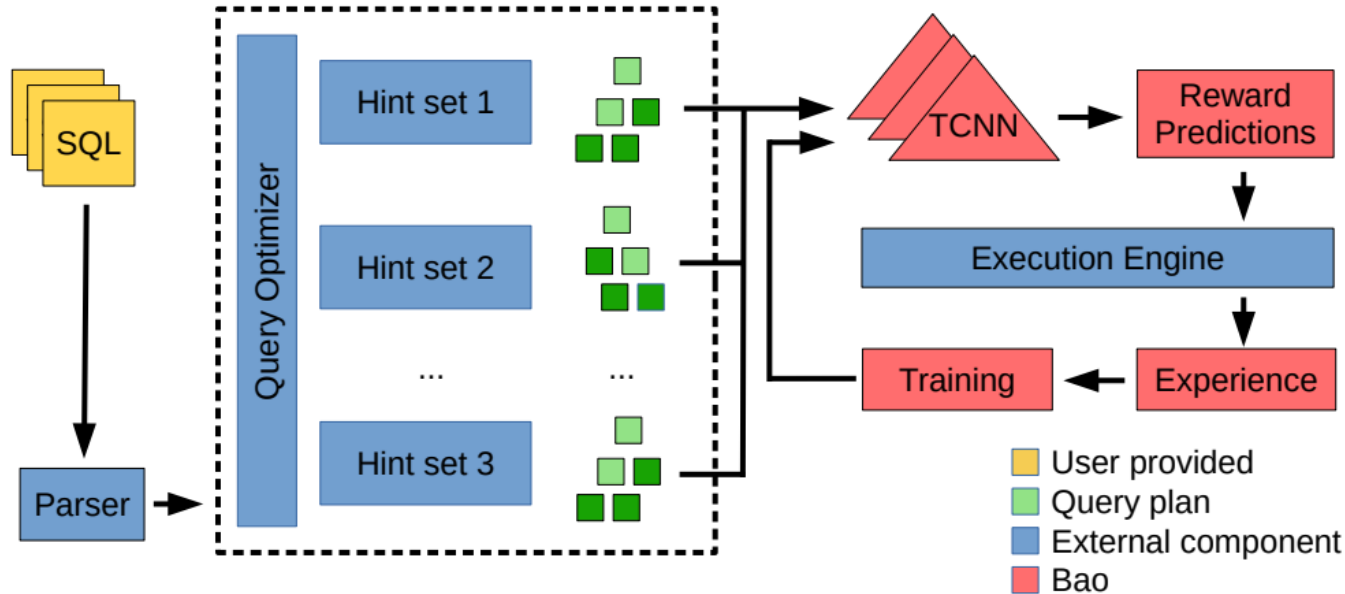
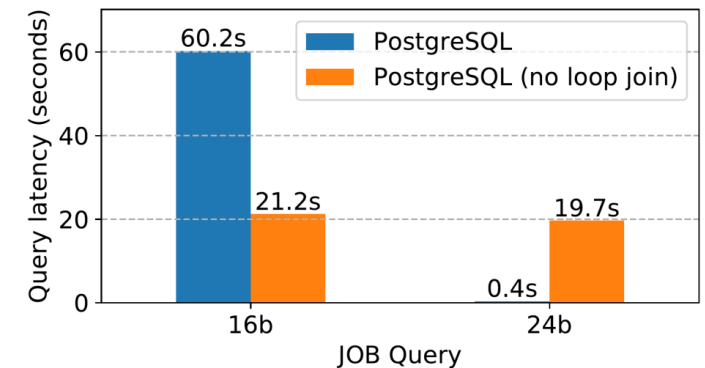
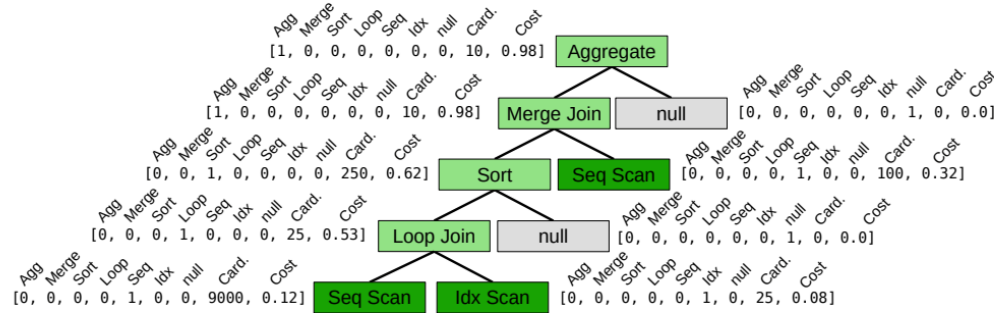


Figure 2: Bao system model

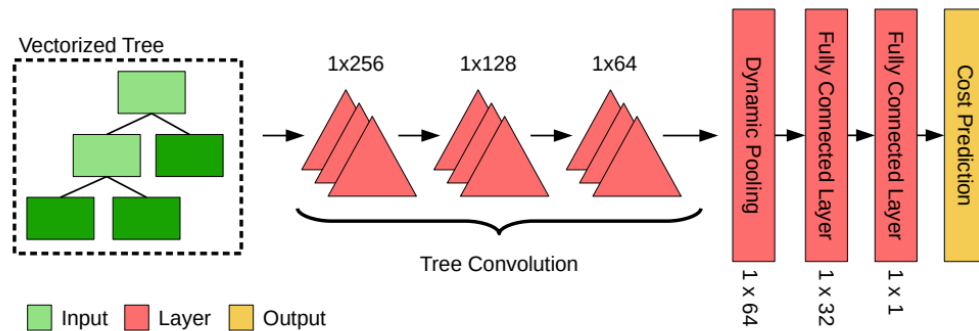
```
enable_hashjoin = false;  
enable_mergejoin = false;  
enable_nestloop = false;  
enable_indexscan = false;  
enable_seqscan = false;  
enable_indexonlyscan = false;
```



# Bao: Making Learned query Optimization Practical (SIGMOD 2021 **Best Paper**)



**Figure 4: Vectorized query plan tree (vector tree)**



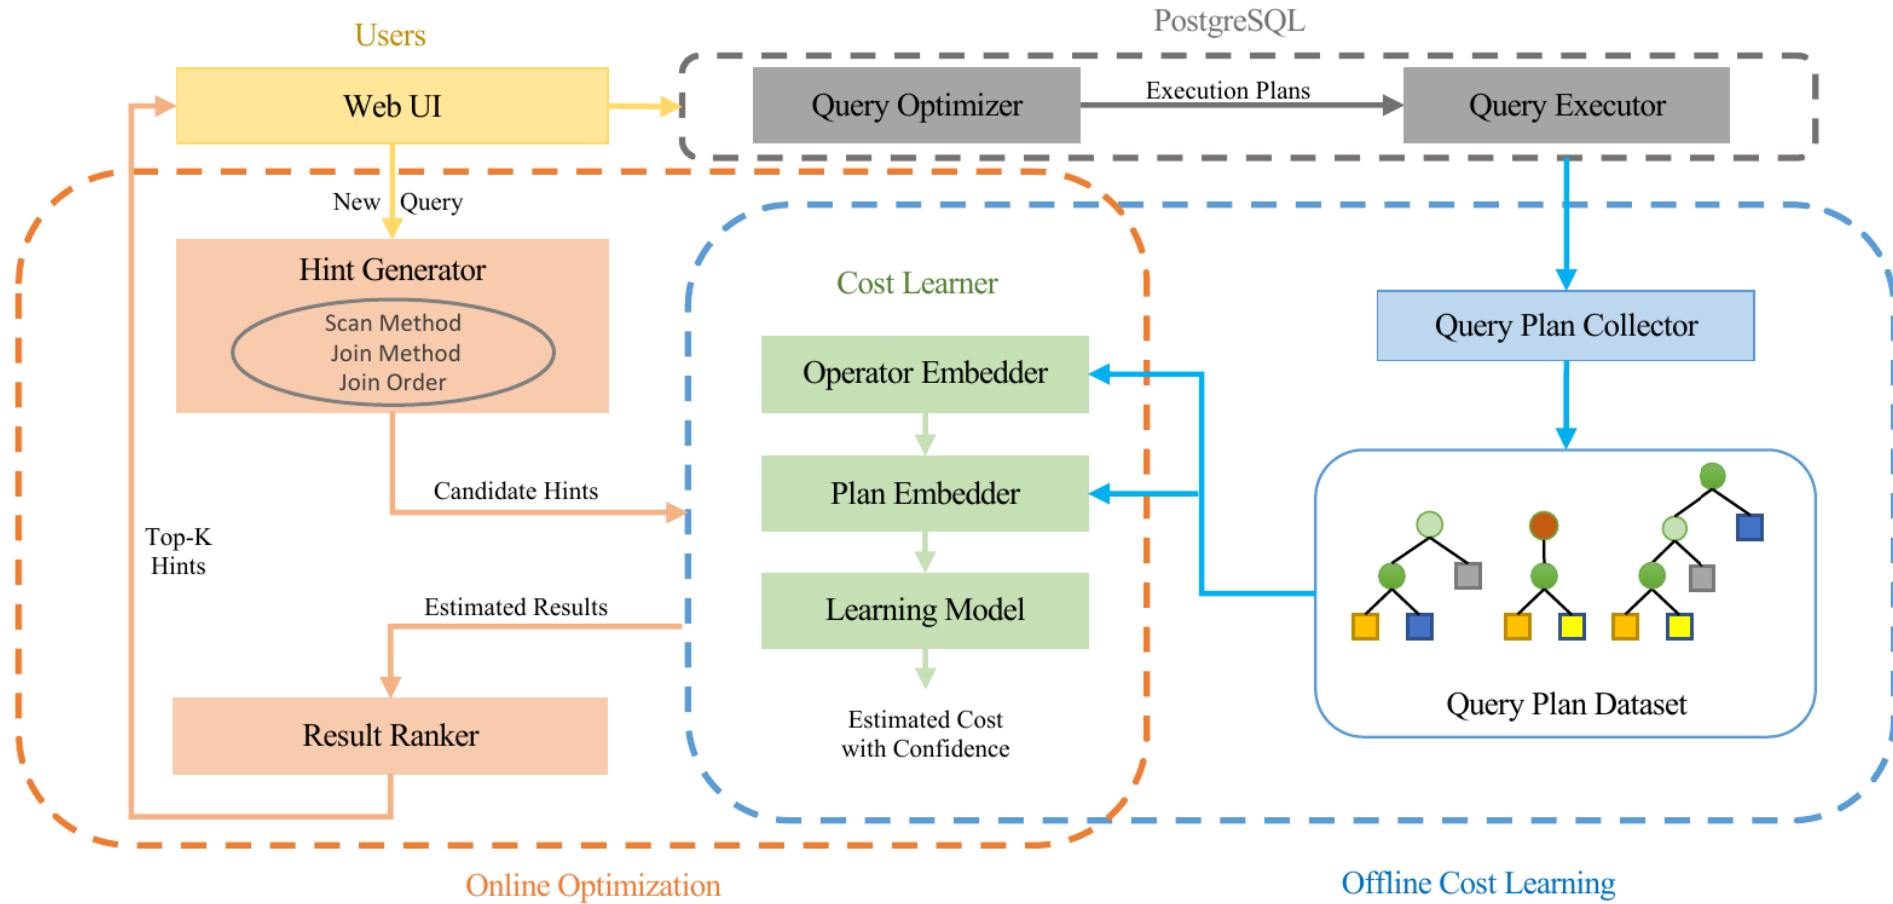
**Figure 5: Bao prediction model architecture**

Q	PG	Bao worst	Bao best
q1	275.415884	12.206382	6.005776
q2	71.049927	198.310226	9.242487
q3	10.982070	290.048801	10.805816
q4	26.890862	26.966064	1.527303
q5	9.692364	9.354480	1.350012
q6	21.741243	19.851484	7.341236
q7	51.935738	51.321676	7.288905
q8	28.725613	15.981973	5.995388
q9	15.645138	16.394102	7.327004
q10	11.720967	9.688347	7.373339
q11	15.163100	7.686548	5.853226
q12	12.934380	9.379889	4.565600
q13	18.687008	11.803825	3.417922
q14	11.100027	14.864732	7.060695
q15	9.641760	8.258874	4.153027
q16	5.312640	7.992982	1.221813
q17	6.404161	17.702658	5.868285
q18	11.912653	20.336241	6.772051
q19	9.943220	33.939818	10.330661
q20	0.143906	0.679753	0.344254
q21	1.022706	1.292618	0.921263
q22	16.113360	51.231996	8.196555

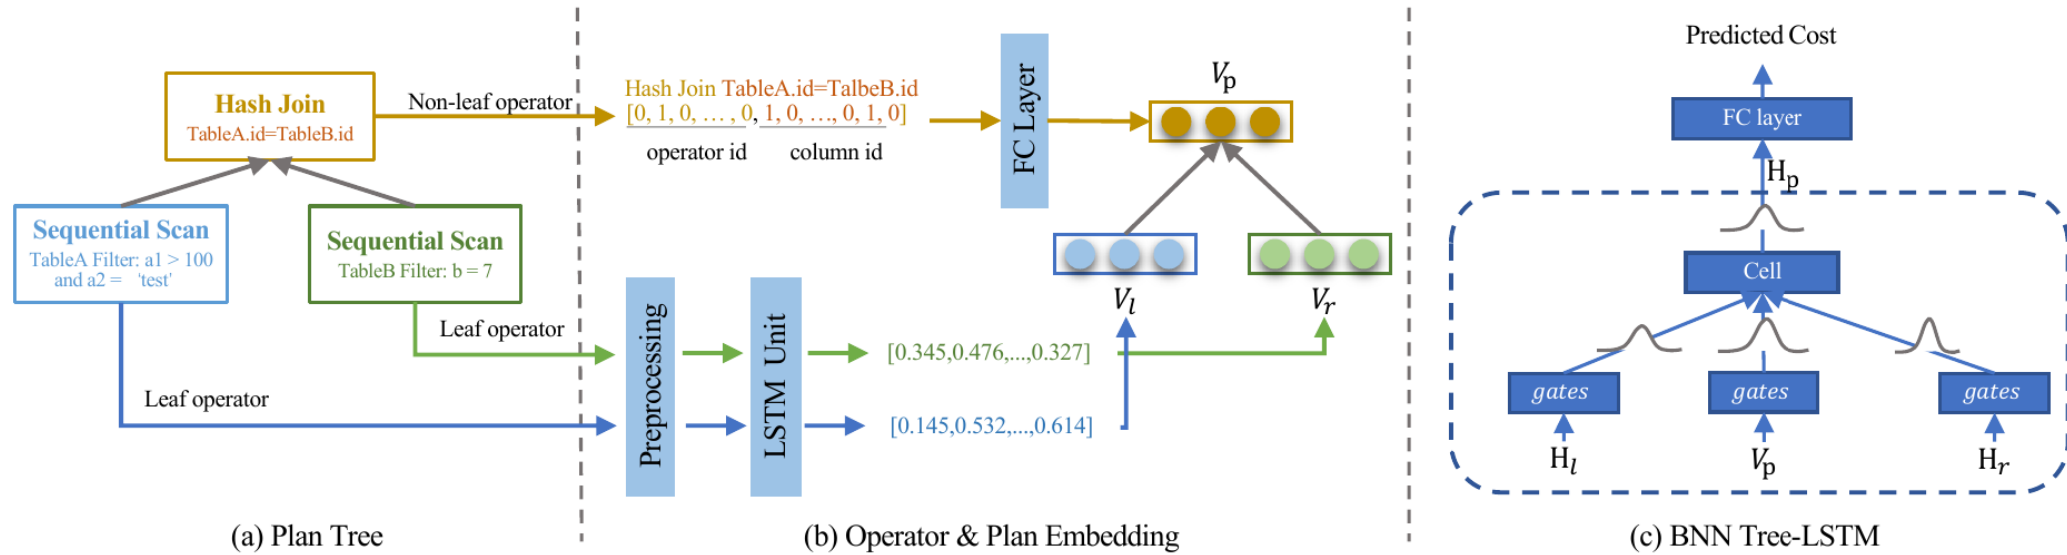
# Deep✱ A Learned Query Optimizer

Luming Sun, Tao Ji, Cuiping Li, Hong Chen  
SIGMOD 2022 Demo Track

# System Architecture



# Cost Learner Intervals



# Contributions

- Practical
  - DeepO is integrated with real DBMS (PostgreSQL).
- Novel
  - DeepO adopts BNN-based Tree-LSTM network.
  - DeepO offers fine-grained query optimization.
- User-friendly
  - DeepO has a web UI for interactive operation.
- Effective
  - DeepO can optimize queries.

# Takeaways

- Reinforcement learning is suitable for the task.
- Join order optimization (query plan optimization) is a very difficult task.
- Challenges and opportunities coexist, as do hardships and hopes.



# Thanks!

[sunluming@ruc.edu.cn](mailto:sunluming@ruc.edu.cn)  
<https://lumingsun.github.io/>